

Threads

Module 07

Reading

This module covers section 2.2 in the text.

Why Multi-process?

Multi-process programming allows use to write “programs” that can do two or more things at once!

Operating in parallel is a powerful model, but process are considered a **heavy-weight** approach.

Heavyweight?

Process have their own:

- Address space
 - Global variables
 - Code (separate copy)
 - State (Program Counter, Registers, Stack)
 - Open Files
 - Accounting
-
- Creating Processes is expensive
 - Communicating between processes is expensive

When are processes a good idea?

- When process creation and communication are *relatively* infrequent.
- When the jobs of the processes are largely independent.
- Multi-process systems can employ different languages, and be moved across networks more easily

Lightweight Alternative


- Often we want to do things in parallel, but don't want to pay the price of multiple processes.
- **Threads:**
 - Shared Address space
 - Shared code
 - Shared files, accounting, etc.
 - Independent State - PC, Registers, Stack

Threads

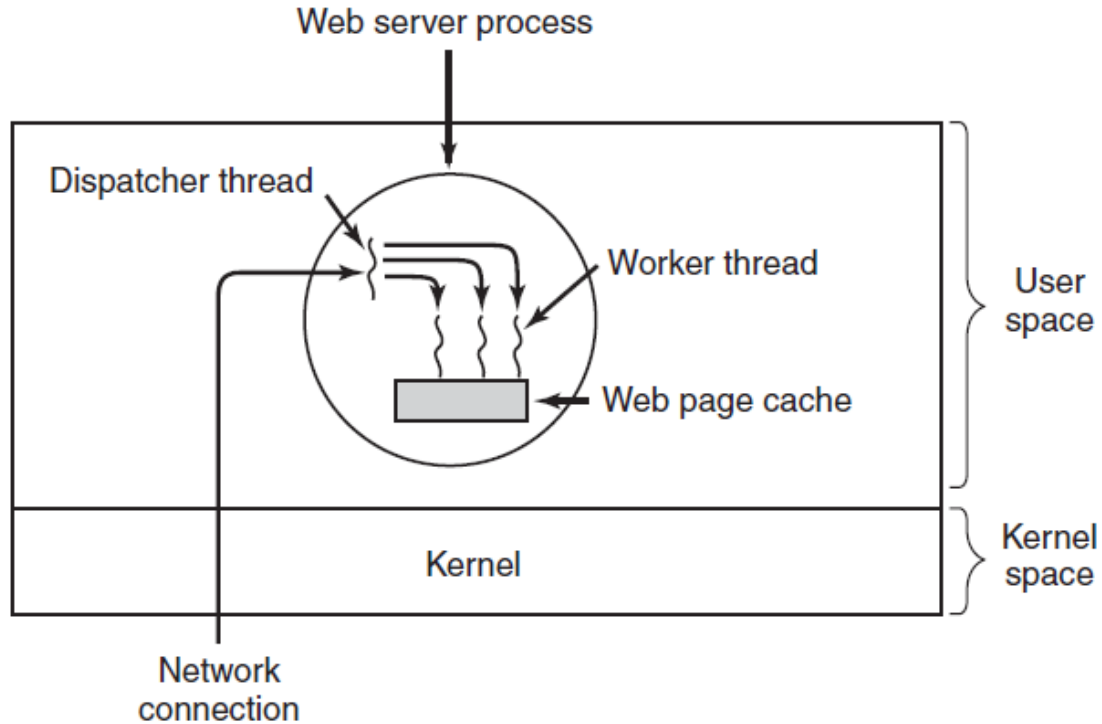
One process can contain many threads

- Each thread has its own execution context - it can be at a different place in the code
- Threads **share global data and the heap** - which facilitates faster (not necessarily easier) communication
- Quick to start

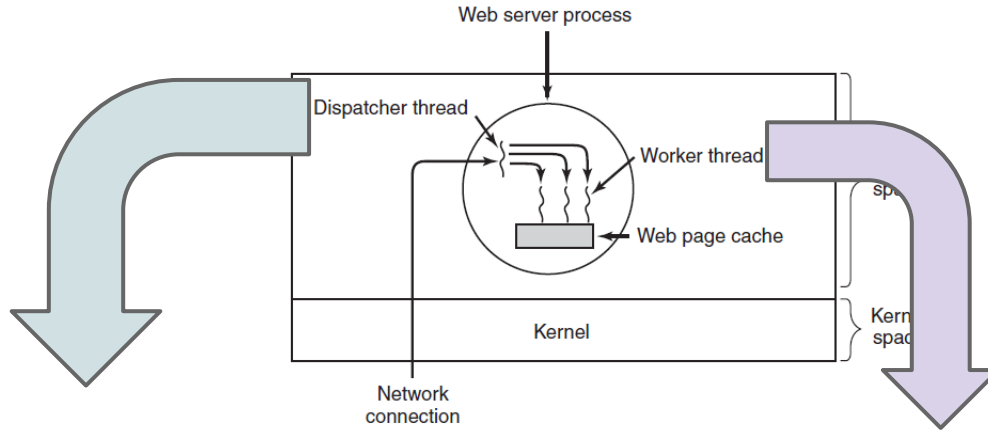
Thread Examples - Word Processor



Thread Example - Web Server



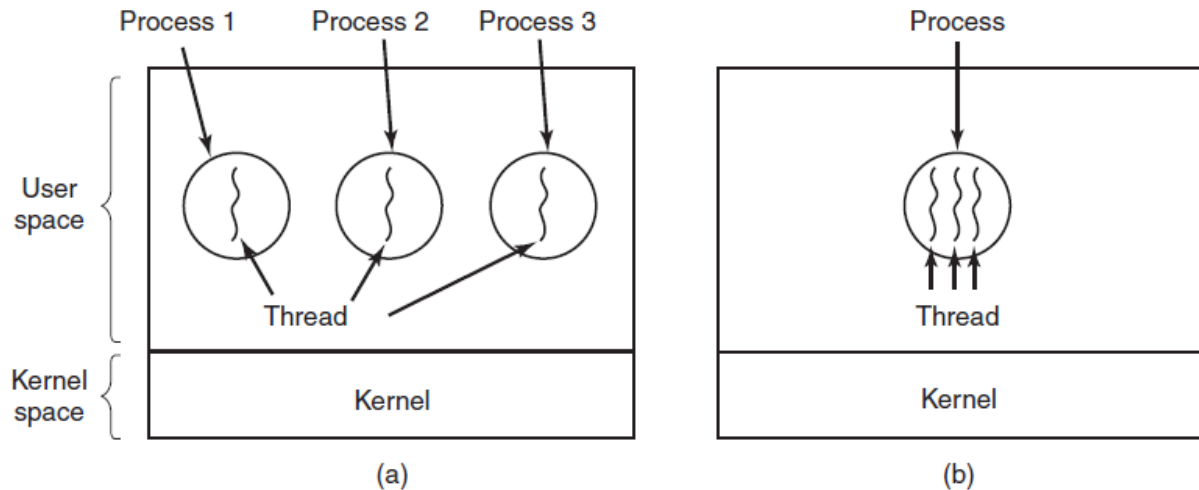
Thread Example - Web Server



```
while (TRUE) {  
    get_next_request(&buf);  
    handoff_work(&buf);  
}
```

```
while (TRUE) {  
    wait_for_work(&buf)  
    look_for_page_in_cache(&buf, &page);  
    if (page_not_in_cache(&page))  
        read_page_from_disk(&buf, &page);  
    return_page(&page);  
}
```

Classical Thread Model



A process *contains* at least one thread


Classical Thread Model

Per process items

- Address space
- Global variables
- Open files
- Child processes
- Pending alarms
- Signals and signal handlers
- Accounting information

Per thread items

- Program counter
- Registers
- Stack
- State



Programming with Threads

- In **POSIX**, there is a set of standardized API functions for managing threads - called **pthread**s
- On **Windows** (and we'll do this now), Win32 API provides a very strong API for threading
- **C++ 11** also created a thread API that allows you to write portable thread code - we'll see this too.
- But first - we'll see that they all require a standard concept - ***function pointers***

“Review” of Function Pointers

In both POSIX and Win32, creating a thread requires you to provide a function pointer

- This defines where the new thread will start running.
- Just like variables, functions have “types”
- Type equivalent to function signature
- The function’s name is a pointer to “code” instead of data

Function Pointers in C

```
// Two function, both with the same signature
```

```
int add(int a, int b);
```

```
int sub(int a, int b);
```

```
void print(int a, int b);
```

```
// Function that takes a and b and performs the operation
```

```
// by executing the function provided:
```

```
int execute(int a, int b, int (*function)(int, int) ) {
```

```
    return function(a, b);
```

```
}
```

```
// a call to the execute function:
```

```
execute (7, 6, sub);
```

```
execute(7, 6, print); // print doesn't have the required signature
```

pthread library

Must include `pthread.h`

Pthreads start in a user defined function with specific signature:

```
void * function_name(void * param)
```

Functions:

- `pthread_attr_init` (set thread attributes to defaults)
- `pthread_create` (create thread)
- `pthread_join(id)` (wait for thread (id) to terminate)

compile with `-lpthread` option

```
g++ -o myprog myprog.cpp -lpthread
```


pthread example

Let's take a look at `simple-thread.c` for an example

Parallel Example

Now lets take a look at simulating parallel work with `posix - thread-summary-posix.cpp`

Note - I'm switching over to C++ for a bit!

Windows!



We've neglected Windows so far! *Poor Windows.*

- POSIX uses integers to represent most operating system resources:
 - pid (process id) from fork
 - fd (file descriptor) fopen, pipe
 - tid (thread id) from pthreads
- Windows uses **HANDLE** data structures

Threads on Windows

HANDLE CreateThread(...)

- security attributes (NULL)
- default Stack Size (0)
- thread function
- parameters to thread function
- creation flags (0)
- &thread identifier (DWORD, not really used)

Thread function has strange signature:

```
DWORD WINAPI functionName(LPVOID parameters)
```

Waiting for Threads to stop

- For pthreads, the term is “join” - we join on a thread which means “wait until thread is finished”
- On Windows, there's a pretty logical set of functions for this:

```
WaitForSingleObject(ThreadHandle, milliseconds)
```

```
WaitForMultipleObjects(count, ThreadHandles, all, milliseconds)
```

Windows Example

Let's rewrite the same summation example, highlighting the Win32 API versus POSIX

`thread-summation-win32.cpp`

C++ 11 Standard



Modern applications almost *always* use threads
- and most languages that have *runtimes*
provide their own API

- Java Threads call corresponding system calls on target platform - OS independent code!
 - Same with C#, Python, Ruby as well...
- Until recently, C++ **didn't** - *and this made people sad.*

C++ 11 Threads

C++ Threads are **objects**.
This is pretty standard in
most OO languages

`std::thread`

The constructor accepts a
function, and immediately
starts a new thread

```
// thread example
#include <iostream>           // std::cout
#include <thread>              // std::thread

void foo()
{
    // do stuff...
}

void bar(int x)
{
    // do stuff...
}

int main()
{
    std::thread first (foo);
    std::thread second (bar,0);
    std::cout << "main, foo and bar now execute concurrently\n";

    // synchronize threads:
    first.join();    // pauses until first finishes
    second.join();   // pauses until second finishes

    std::cout << "foo and bar completed.\n";

    return 0;
}
```




C++ 11 Threads

Lets look at `cpp11threads.cpp`

You *probably* need to tell your g++ compiler to use C++ 11 features - this will eventually go away

```
g++ -o test cpp11threads.cpp -std=c++11
```

Kernel vs. User Threads



Kernel vs. User Threads

- Kernel threads are visible to the OS, independently schedulable
 - When one thread is blocked, other threads *in same process* can still be run
 - Multiple threads within same process can run on multiple CPU's
- User threads are invisible to the OS
 - One user thread blocks, all the process's user threads are blocked.
 - Cannot run user threads on multiple CPU's

Kernel vs. User Threads

- User threads are fast, but do not offer the same degree of parallelism as kernel threads
- User threads provide the programming model - but not really the performance benefits
- Nearly all modern operating systems support Kernel Threads

Next up...

After the exam, we'll start with Chapter 2.3 - Interprocess (and Inter-Thread) communication

This will be **the most** important portion of the semester for your programming skills.