

# Scheduling

## Module 09

# Reading

This module covers Chapter 2.4 in the text.

# Man behind the curtain...

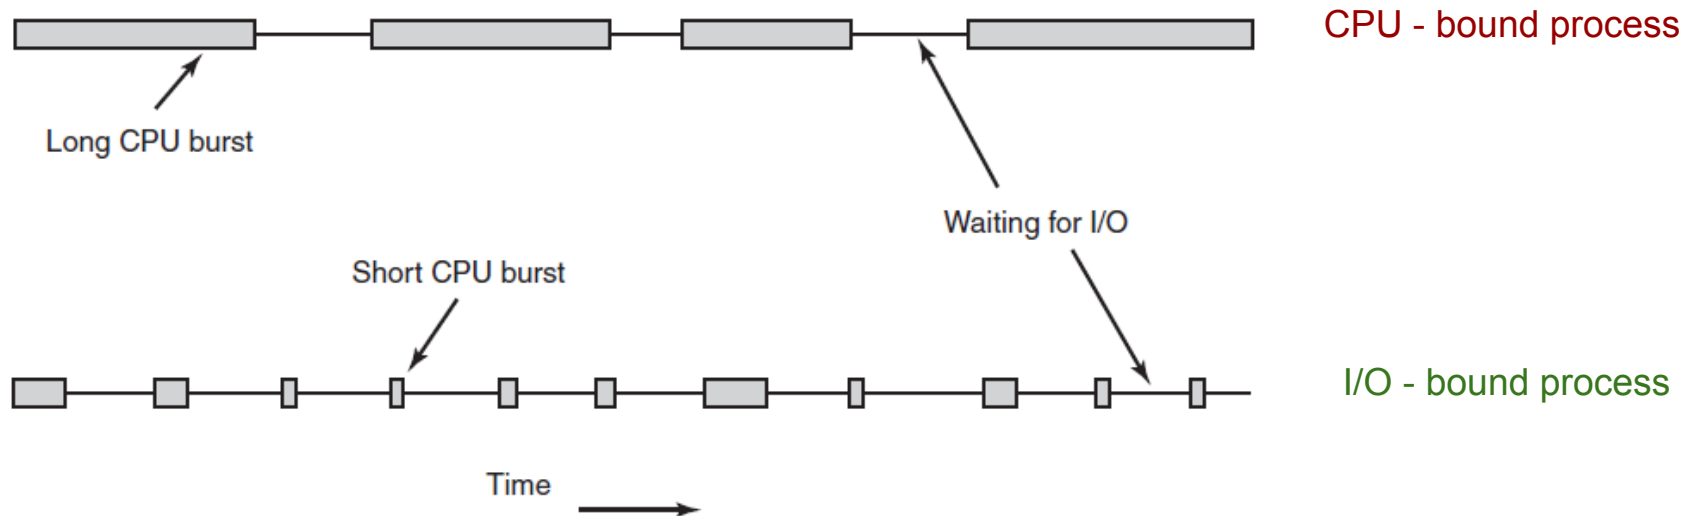
We've been writing synchronization code in a state of paranoia - because we never know when the schedule is going to interrupt us!



Now we will look at how the schedule actually makes these decisions, and how deterministic it really is.

# What are we scheduling?

All processes alternate between CPU instructions and an I/O call



# Categorizing processes

**I/O Bound** programs don't compute much between I/O calls - little CPU usage

**CPU Bound** programs have lots of calculations between I/O calls

*The length of the I/O is irrelevant*

# What are we actually scheduling?

- We are scheduling CPU-bursts.
- Once an I/O call is made, the process is blocked and we schedule another **ready** process's CPU burst

# When does the scheduler schedule?

- Its critical to understand that the scheduler can only run when the OS is on the CPU
- It must make a decision who should run next, and set the PC appropriately.
- Once the PC is set - the scheduler is ***not running***.

# When can the scheduler run?

- Once a process starts running, the OS won't run until one of the following occurs:
  - The process makes a system call like **fork** or **exit**
  - The process yields voluntarily (**sleep**, **yield**)
  - The process issues a blocking call (**read**)
  - An interrupt occurs (pending I/O completes)

```
while ( true );    // oh... no.
```



# Preemptive Scheduling

- **Nonpreemptive** systems are at the mercy of processes triggering a scheduling opportunity.
- **Preemptive** systems employ a **hardware timer** to issues an **interrupt** after a time period.
  - This time period is called a *quantum*
  - The schedule can set the timer before starting a process.

# When can the scheduler run?

- Once a process starts running, the OS won't run until one of the following occurs:
  - The process makes a system call like **fork** or **exit**
  - The process yields voluntarily (**sleep**, **yield**)
  - The process issues a blocking call (**read**)
  - An interrupt occurs (pending I/O completes)
  - A timer interrupt goes off (quantum expiration)

# Measuring effectiveness

- A scheduling **algorithm** is a procedure for deciding which process runs next
- If we are going to *evaluate* scheduling algorithms, we need **metrics**
  - CPU Utilization
  - Throughput
  - Turnaround Time
  - Wait Time
  - Response Time
  - Fairness

Different types of systems need to optimize different metrics

# Measuring effectiveness

- A scheduling **algorithm** is a procedure for deciding which process runs next
- If we are going to *evaluate* scheduling algorithms, we need **metrics**
  - CPU Utilization
  - Throughput
  - Turnaround Time
  - Wait Time
  - Response Time
  - Fairness

Different types of systems need to optimize different metrics

← We're going to defer "priority" for a moment...

# Types of Systems

**Batch Systems:** Payroll systems, banking systems, data collection tasks..

**Interactive Systems:** Likely every computer you've ever used.

**Real-Time Systems:** Missile guidance, satellite control, and other silly programs.

# Batch System Priorities

Maximize Throughput

Minimize Turnaround



These are clearly linked - but not exactly the same.

Throughput can be achieved by ignoring long jobs!

Keep CPU utilization high



Drives Throughput/Turnaround

The CPU should never be idle  
**Minimize time OS is on CPU**


# Interactive Systems

- When users are clicking and typing, priorities change.
  - Users was **their program** to run **now**.
  - Overall system performance is not quite as important as **foreground** process performance
  - Ability to switch tasks is **critical**




# Interactive Systems

- Interactive Systems should minimize **wait time** and minimize **response time**.
- As we will see, response time can only be minimized at the expense of CPU utilization... **and we know what that means!**





# Real-time Systems



- Real-time does not mean “super fast”
- Real time means you have very specific times when certain things need to execute
  - It might be every ten minutes, run a 10 second task -  
the point is we don't mean “around” every ten minutes!
- Real-time systems are another animal, we're going to leave them alone for now

# Scheduler Implementation

- Processes are modeled by entries in a **process table / process control block**
  - Entries contain *accounting information*, along with other information about *restarting* the process
  - The schedule maintains a **ready queue** of all processes ready to run.
- When the schedule runs, it consults the table and its own records and picks a process
  - It restores registers/PC/stack etc of the chosen process.
  - This is called a **context switch**.

# First Come, First Serve

Similar to FIFO queue, this algorithm simply picks the *next* process from the **ready queue**.

- **This is a non-preemptive algorithm:**
  - The process selected runs until it blocks, yields, or terminates

Let's model a hypothetical execution of a set of CPU bursts with FCFS scheduling

# Shortest Job First (SFJ)

Is the wait time for FCFS any good? What would be optimal?

- Select the shortest job first
  - Still non-preemptive

Let's model a hypothetical execution of a set of CPU bursts with SJF scheduling

# Why is SJF Optimal?

Let's assume we have four bursts to schedule (A, B, C, and D), run in order

A waits no one  $\emptyset$

B waits for A A

C waits for A and B  $A + B$

D waits for A, B, and C  $A + B + C$

---

Total Wait Time:  $3A + 2B + C$

# What's the shortest job?

- A short job is really a **short CPU burst**
  - Processes with short CPU burst are I/O Bound

## Critical Insight:



**Prioritize I/O bound processes over CPU bound processes!**

- This also makes sense from a resource optimization perspective!

# Preemptive SJF

- We can also simulate preemption, and a dynamic ready queue, and still employ SJF
  - Model *arrival* times along with CPU bursts
  - Schedule updates burst time when process is interrupted.
  - Remember - new arrivals will constantly enter the ready queue (I/O completions)

Let's model a hypothetical execution of a set of CPU bursts with “Least Remaining” scheduling

# Implementing SJF?



```
while (true) {  
    x = random()  
    if ( x % 2 == 0 ) {  
        read_disk();  I/O (CPU Burst is over)  
    }  
    else {  
        do_math();  CPU burst's reign of terror continues.  
    }  
}
```



# Approximating Shortest Job First

- We could keep accounting data and use the past to predict the future
- Need to be careful, the more time we spend doing this, the worse CPU utilization gets!

Later we'll see how to do this better.

# Lets look at FCFS, with preemption

- FCFS is great for batch systems, but bad for interactive - poor task switching
- Round Robin employs FCFS, but uses a **quantum** to preempt running process
- What happens to response time as quantum time changes?
- What happens to CPU utilization?

Let's model a hypothetical execution of a set of CPU bursts with Round-Robin scheduling

# Priorities

- We've established that I/O should be given priority over CPU - bound processes
  - I like to call this *dynamic* or *internal* priority
  - It can change as we get a sense of how I/O or CPU-bound the process is.
- However - processes may be **assigned** priorities
  - I'll call this *static* or *external* priority

# Priority Scheduling

We can simply always choose the highest priority process

Let's model a hypothetical execution of a set of CPU bursts with **Priority** scheduling

**Problem:** Starvation

**Solution:** Aging

# Lottery Scheduling


- Priority Scheduling is very deterministic:
  - The low priority process **never** runs!
- Lottery scheduling provides a **probabilistic** result, where high priority task are simply *likely* to run.

# Scheduling multi-core systems

- **Goal:** Load balancing
  - We don't want one core to be idle while a bunch of processes are waiting on another!
- **Implementation Options**
  - Split-Ready Queue
  - Common-Ready Queue

# Processor Affinity

Common-ready queue seems to make sense for load balancing... but wait!



- The desire to keep a process on a specific CPU is called processor **affinity**.
- In practice, this is so important that split ready queues are typically used instead of a common-ready queue.
- **Pull Migration**: If a CPU's ready queue is empty, go look for another ready queue to pull from.

# Real World Example: Sun Solaris

- Priority - Based Scheduling
  - Separate Queue for each priority level:
  - Real-time, System, Time Sharing, Interactive
- Time-sharing and Interactive have ***dynamically changing priorities***
  - 0-60, 0 being the highest



# Real World Example: Sun Solaris

			Priority after...	
	Priority	Quantum Time	Quantum Expiration	Sleep/IO Completion
Lowest	0	200	0	50
	10	160	0	51
	20	120	15	52
	35	80	25	54
	50	40	40	58
Highest	59	20	49	59

Why penalize quantum completion?

Why boost on I/O completion?

# Real World Example: Windows

## Priority Scheduling via a “grid”

	real-time	high	above normal	normal	below normal	idle
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

# Real World Example: Windows

Columns are *externally defined through Win32 API - static*

	real-time	high	above normal	normal	below normal	idle
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

# Real World Example: Windows

Rows are dynamically adjusted

	real-time	high	above normal	normal	below normal	idle
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

# Real World Example: Linux

- Priority based again!
- Priorities are externally defined (only)
  - High priority tasks get larger quantum
  - Processes occupy a **split ready queue**.
    - Expired
    - Active

# Real World Example: Linux

“Ready Queue”



- Only process in Active queue are selected
- On pre-emption, kicked into Expired queue
- When active empties, switch pointers

# Take-away from Real-World Examples

- Real Operating systems use a **blend** of strategies
  - They all allow external priorities
  - They all favor I/O and penalize CPU
  - They achieve their goals differently!

## Next Up...

After Exam 2 we will start memory management

Chapter 3.1-3.2